

# Theyaneshwaran Jayaprakash

+1 (812) 822-7804 | [theyaneshj17@gmail.com](mailto:theyaneshj17@gmail.com) | [linkedin.com/in/theyaneshj17](https://www.linkedin.com/in/theyaneshj17) | [github.com/theyaneshj17/](https://github.com/theyaneshj17/)

## Education

### Indiana University Bloomington

Master of Science, Data Science; GPA: 4.00

August 2024 - May 2026

Bloomington, United States

### College of Engineering Guindy, Anna University

Bachelor of Engineering, Electronics and Communication Engineering; GPA 3.70

August 2016 - April 2020

Chennai, India

## Research Experience

### ML Research Assistant

IU School of Medicine

September 2024 - Present

Indiana, USA

- Developed a self-supervised autoencoder in PyTorch, guided by Gaussian Mixture Model clustering for classification tasks, achieving **92% accuracy in Alzheimer's subtype discovery** ([First Author, AAIC 2025](#); Fellowship Recipient).
- Validated identified subtypes via multidimensional analysis across neuroimaging and behavioral data domains, employing ANOVA, hypothesis testing, and **Bayesian MCMC inference** to quantify statistical significance.

## Professional Experience

### AI Engineer Intern

FocusKPI, Inc

June 2025 - Present

California, USA

- Architected and deployed a scalable **Model Context Protocol (MCP) client-server system** on AWS ECS, processing 75,000+ FDA medical device records with RAG and BGE embeddings for semantic search across 200 API fields.
- Engineered a real-time **multi-modal AI agent** with desktop overlay integration, achieving sub-5 second latency for speech transcription and visual analysis using Whisper and GPT-4V, to generate automated summaries and recommendations.
- Orchestrated **multi-agent workflows** using LangChain and LangGraph for context management across audio, visual, and text data streams.

### Senior Data Engineer

UnitedHealth Group

March 2023 - August 2024

Chennai, India

- Optimized ELT efficiency by parallelizing workflows and **tuning SQL performance**, reducing monthly processing time from **5 to 2 days** and cutting compute costs by 50% (6K to 3K Data Warehouse Units).
- Migrated **50TB+ data** from Synapse to Snowflake using PySpark and SQL with ADF for phased migration with continuous sync, implementing **dbt transformation** layer for analytics.
- Developed **Power BI semantic models** in Microsoft Fabric, optimizing Direct Lake/Query modes with row-level security. Built 10+ interactive dashboards, **enhancing performance by 67%** via fine-tuned DAX queries.

### Data Engineer

UnitedHealth Group

October 2020 - February 2023

Hyderabad, India

- Integrated **25+ data sources** (Parquet, JSON, CSV, Delta) from Oracle, Teradata, and Databricks into Azure Synapse using a **Medallion architecture**, improving ELT scalability.
- Built a FHIR streaming pipeline on Databricks with Delta Lake and PySpark, processing **5,000+ records/min** from Azure Event Hubs using optimized micro-batch triggers and partitioning.
- Migrated APIs from Azure VMs to Container Instances, enabling cross-cloud Snowflake queries, **reducing costs by 60% (\$30K annually)** and improving execution time from 60s to 15s using C# and SQL optimizations.
- Automated CI/CD for 7 repositories with GitHub Actions, **halving release cycles** and improving deployment reliability.

## Skills

**Programming & Technologies:** Python, R, C++, SQL, Git, Docker, Kubernetes, Hadoop, PySpark, A/B Testing

**Data Engineering & Cloud:** Azure (Synapse, ADF), AWS (S3, Lambda, Redshift), Databricks, Snowflake, dbt, Airflow

**Data Analysis & Visualization:** Pandas, NumPy, Matplotlib, Tableau, Power BI, ggplot2, dplyr

**Machine Learning & AI:** PyTorch, TensorFlow, Scikit-learn, MLflow, RAG, Transformers, Causal Inference

## Projects

### IntelliFlow Enterprise AI Orchestration System ([GitHub](#)) | LLM, Neo4j, FastAPI, MCP

August 2025

- Architected LLM-powered workflow orchestration with intent classification across 4 MCP servers, constructing Neo4j knowledge graph with **50K+ entities** enabling **multi-hop reasoning** for IT, data, and research automation.
- Developed secure **text-to-SQL engine** with parameterized query validation, real-time ServiceNow integration, and **GraphRAG**-based research capabilities for enterprise query processing.